

Ajeet K. Jain, M. Narsimlu
(ML TEAM)- SONET, KMIT, Hyderabad

Session – 31

This session deals with

Introduction to Case Study



Problem statement

- Subsidy Inc. delivers subsidies to individuals based on their income
- Accurate income data is one of the hardest piece of data to obtain across the world
- Subsidy Inc. has obtained a large data set of authenticated data on individual income, demographic parameters, and a few financial parameters
- Subsidy Inc. wishes us to :
Develop an income classifier system for individuals

The Objective is to:

Simplify the data system by reducing the number of variables to be studied, without sacrificing too much of accuracy. Such a system would help Subsidy Inc. in planning subsidy outlay, monitoring and preventing misuse.



```
#To visualize the data
import seaborn as sns
#To work with dataframes
import pandas as pd
#To perform numerical operations
import numpy as np
#To partition the data
from sklearn.model_selection import train_test_split
#importing the library for logistic regression
from sklearn.linear_model import LogisticRegression
#importing performance metrics
from sklearn.metrics import accuracy_score, confusion_matrix
```



```
#importing data
data_income=pd.read_csv("income.csv")
#create a copy of original data
df_income=data_income.copy()
print(df_income.describe())
```

	age	capitalgain	capitalloss	hoursperweek
count	31978.000000	31978.000000	31978.000000	31978.000000
mean	38.579023	1064.360623	86.739352	40.417850
std	13.662085	7298.596271	401.594301	12.345285
min	17.000000	0.000000	0.000000	1.000000
25%	28.000000	0.000000	0.000000	40.000000
50%	37.000000	0.000000	0.000000	40.000000
75%	48.000000	0.000000	0.000000	45.000000
max	90.000000	99999.000000	4356.000000	99.000000




```
#importing data  
data_income=pd.read_csv("income.csv")  
#create a copy of original data  
df_income=data_income.copy()  
print(df_income.info())
```



```
RangeIndex: 31978 entries, 0 to 31977
Data columns (total 13 columns):
age                31978 non-null int64
JobType            31978 non-null object
EdType             31978 non-null object
maritalstatus      31978 non-null object
occupation         31978 non-null object
relationship       31978 non-null object
race               31978 non-null object
gender             31978 non-null object
capitalgain        31978 non-null int64
capitalloss        31978 non-null int64
hoursperweek       31978 non-null int64
nativecountry      31978 non-null object
SalStat           31978 non-null object
```




Data analytics framework



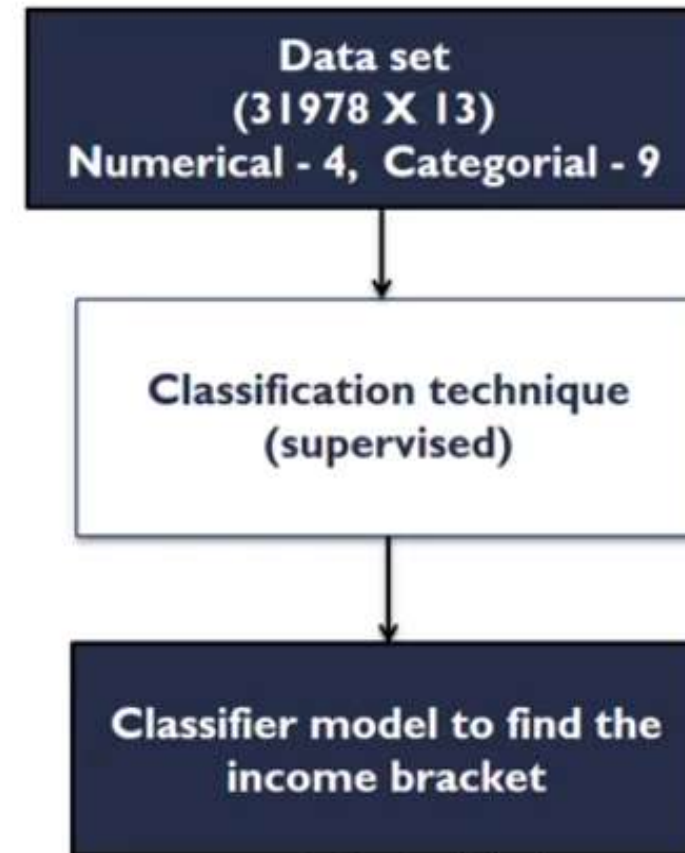


Framework

- Problem conceptualization
 - Develop an income classifier for individuals with reduced no. of variables
- Problem characterization- Classification

Apriori Known:

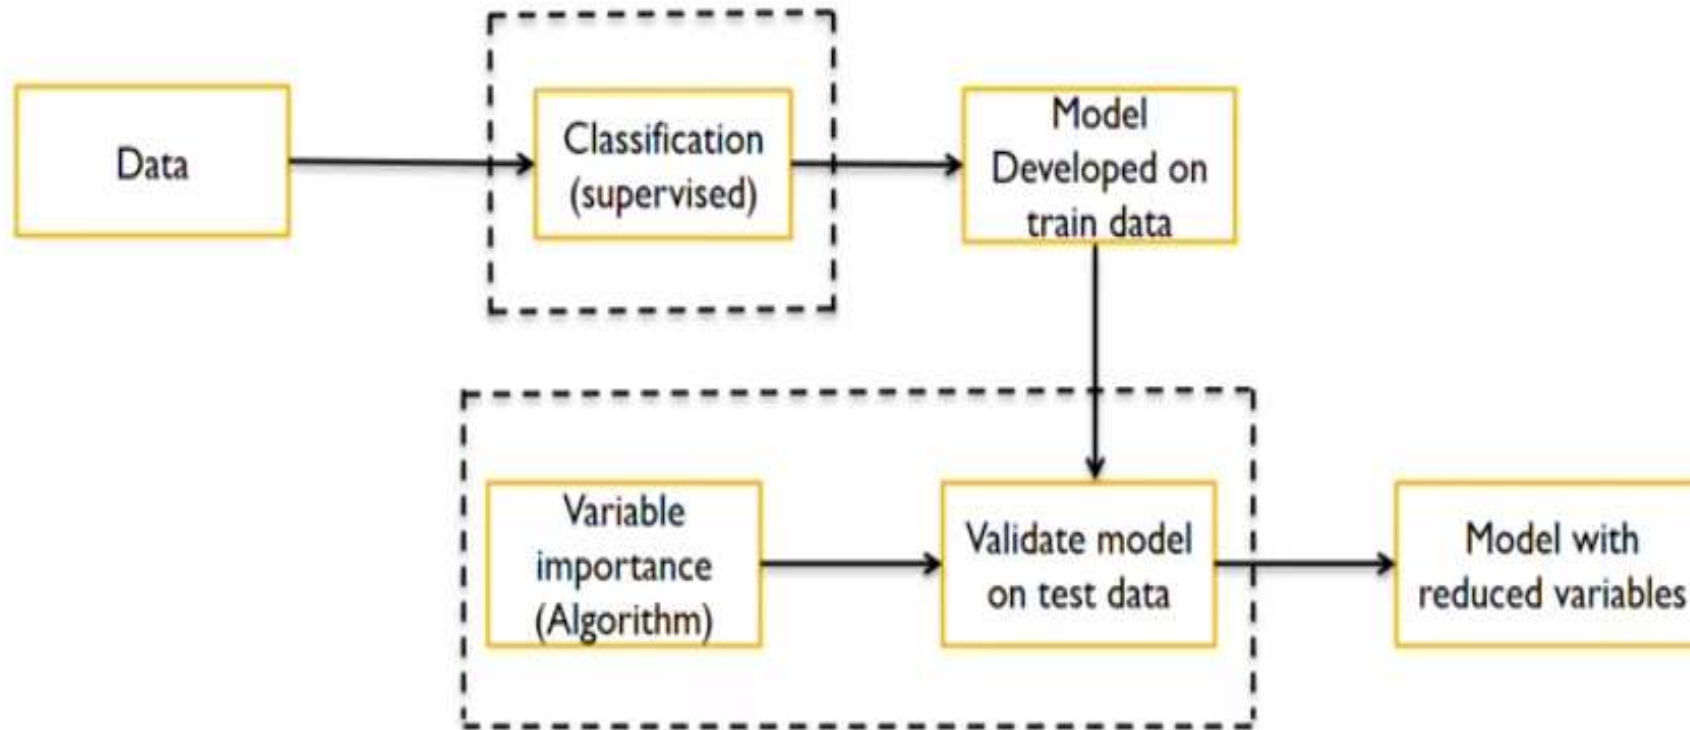
- ✓ Dependent variable
categorical (binary)
- ✓ Independent variables
numerical + categorical





Framework

- Flow chart:





Framework

- Solution conceptualization
 - Identify if data is clean
 - Look for missing values
 - Identify variables influencing salary status and look for possible relationships between variables
 - Correlation, chi-square test, box plots, scatter plots etc.
 - Identify if categories can be combined
 - Build a model with reduced number of variables to classify the individual's salary status to plan subsidy outlay, monitor and prevent misuse



Framework

- Method identification
 - Logistic Regression
 - Random Forest
 - K Nearest Neighbors
- Realization of solution
 - Evaluate performance metrics
 - If assumptions are satisfied and solutions are acceptable then model is



```

"""
EDA
1.getting to know the data
2.Data Preprocessin
3.Cross tables and data visualization
"""

#1.Getting to know the data
#To check variables data type
print(df_income.info())
#To find the missing values in each feature
print(df_income.isnull().sum())
#No missing values
#Summary of numerical variables
print(df_income.describe())
#Summary of categorical variables
print(df_income.describe(include="O"))
#Frequency of each categories
print(df_income["JobType"].value_counts())

```



```
#importing data
df_income=pd.read_csv("income.csv")
#checking for unique classes
print(np.unique((df_income["JobType"])))
print(np.unique(df_income["occupation"]))
#checking other special characters in the dataset
data=pd.read_csv("income.csv",na_values=[" ?"])
#Check missing values in each feature
print(data.isnull().sum())
missing=data[data.isnull().any(axis=1)]
```

#To consider one missing column
print(missing)



```
Private                22286
Self-emp-not-inc      2499
Local-gov             2067
?                     1809
State-gov             1279
Self-emp-inc          1074
Federal-gov           943
Without-pay           14
Never-worked           7
Name: JobType, dtype: int64
```

```
Prof-specialty        4038
Craft-repair          4030
Exec-managerial       3992
Adm-clerical          3721
Sales                 3584
Other-service         3212
Machine-op-inspct     1966
?                     1816
Transport-moving      1572
Handlers-cleaners     1350
Farming-fishing       989
Tech-support          912
Protective-serv       644
Priv-house-serv       143
Armed-Forces          9
```



```
#importing data
data=pd.read_csv("income.csv",na_values=[" ?"])
"""
1.missing values in jobtype -      1809
2.missing values in occupation -   1816
3.There are 1809 rows where tow soecific columns
i.e oocupation and jobtype have missing values
4.(1816-1809)=7 => still we have occupation unfilled for these rows.
because,jobtype is "never worked"
"""

data2=data.dropna(axis=0)
#correlation between independent varaibles
corr_rel=data2.corr()
print(corr_rel)
|
```



Week4')

	age	capitalgain	capitalloss	hoursperweek
age	1.000000	0.080154	0.060165	0.101599
capitalgain	0.080154	1.000000	-0.032229	0.080432
capitalloss	0.060165	-0.032229	1.000000	0.052417
hoursperweek	0.101599	0.080432	0.052417	1.000000



```
#importing data
data=pd.read_csv("income.csv",na_values=[" ?"])
data2=data.dropna(axis=0)
#cross tables and data visualization
#extract the columns
print(data2.columns)
gender_tab=pd.crosstab(index=data2["gender"],columns="counts",
                        normalize=True)

print(gender_tab)

#Relation between Gender vs salary
gender_salstat=pd.crosstab(index=data2["gender"],
                           columns=data2["SalStat"],
                           margins=True,normalize="index")

print(gender_salstat)
```



```
Index(['age', 'JobType', 'EdType', 'maritalstatus', 'occupation',
      'relationship', 'race', 'gender', 'capitalgain', 'capitalloss',
      'hoursperweek', 'nativecountry', 'SalStat'],
      dtype='object')
```

col_0	counts
-------	--------

gender	
--------	--

Female	0.324315
--------	----------

Male	0.675685
------	----------

SalStat	greater than 50,000	less than or equal to 50,000
---------	---------------------	------------------------------

gender		
--------	--	--


Female	0.113678	0.886322
--------	----------	----------

Male	0.313837	0.686163
------	----------	----------

All	0.248922	0.751078
-----	----------	----------




```
#importing data
data=pd.read_csv("income.csv",na_values=[" ?"])
data2=data.dropna(axis=0)
#cross tables and data visualization
#frequency distribution of salary status
SalStat=sns.countplot(data2["SalStat"])
"""
75% of people's salary status is <=50,000
25% of people's salary status is >50,000
"""
```



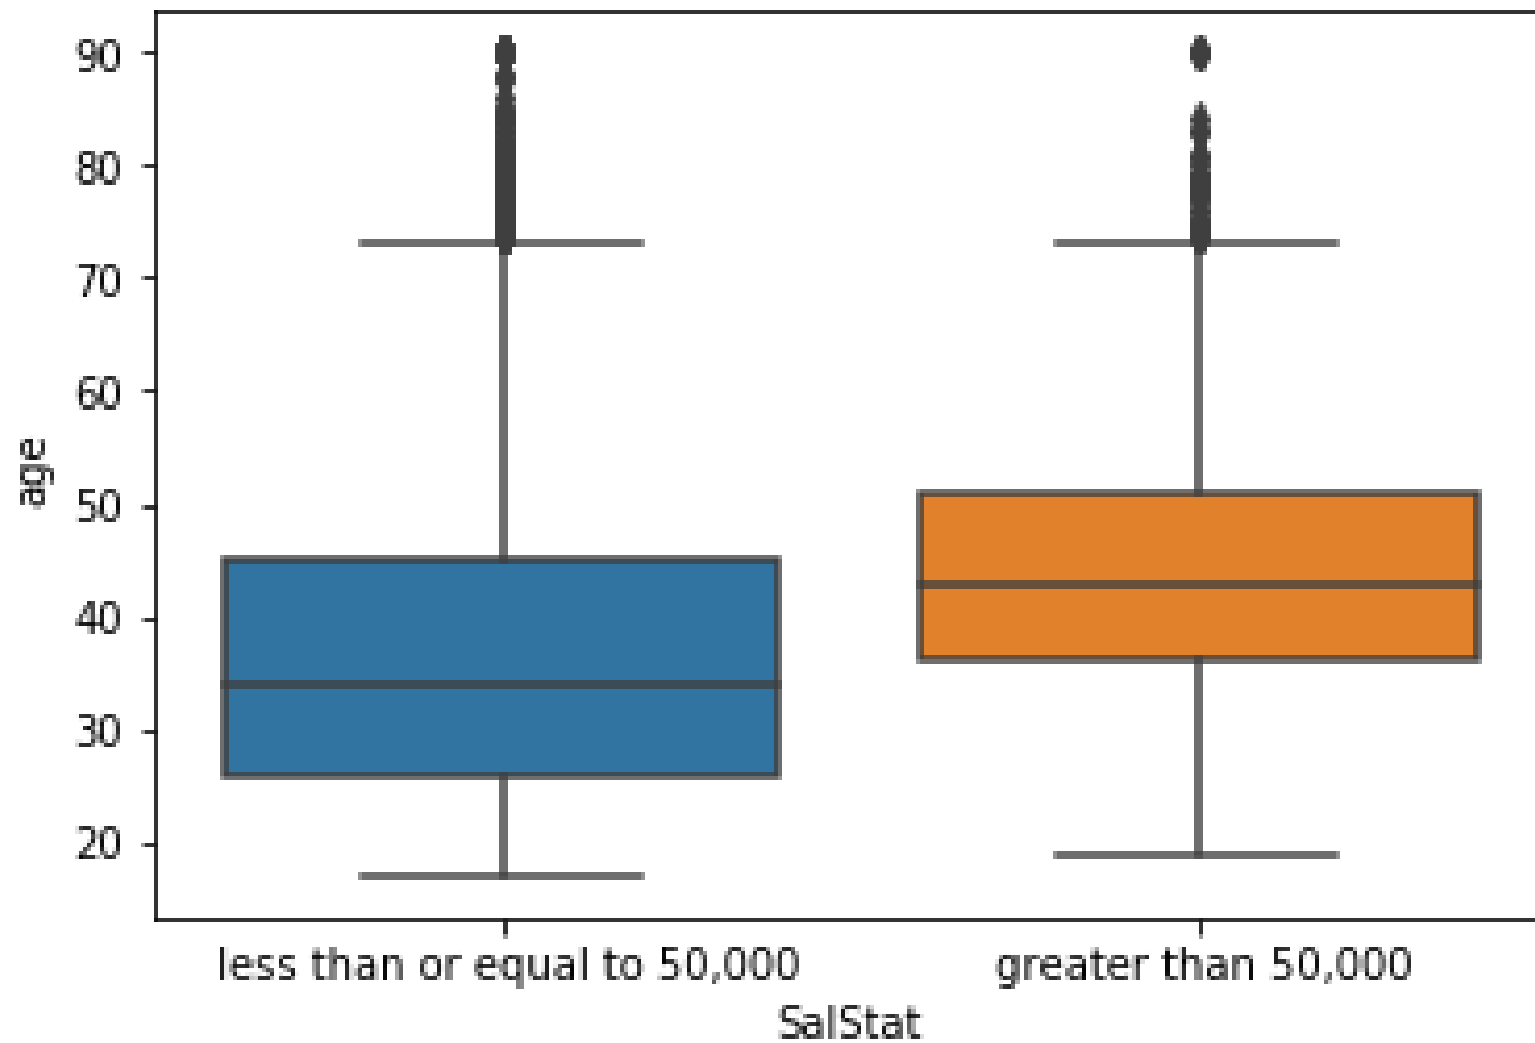


```
#importing data
data=pd.read_csv("income.csv",na_values=[" ?"])
data2=data.dropna(axis=0)
#histogram of Age
sns.distplot(data2["age"],bins=10,kde=False)
#people with age 20-45 age are high in frequency
```





```
#importing data
data=pd.read_csv("income.csv",na_values=[" ?"])
data2=data.dropna(axis=0)
#boxplot Age vs salary status
sns.boxplot("SalStat","age",data=data2)
"""
people with 35-50 age are more likely to earn >50000
people with 25-35 age are more likely to earn <=50000
"""
```





EDA-

1.Jobtype VS salary status


2.create a cross table Jobtype Vs salary status

```
salStat_tab=pd.crosstab(index=data2["JobType"],  
                        columns=data2["SalStat"],  
                        margins=True,normalize="index")  
print(salStat_tab)  
#sns.countplot(y="JobType",data=data2,hue="SalStat")  
"""
```

```
The above table 56% of self employed  
people earn more than 50000 USD per year  
"""
```



SalStat	greater than 50,000	less than or equal to 50,000
JobType		
Federal-gov	0.387063	0.612937
Local-gov	0.294630	0.705370
Private	0.218792	0.781208
Self-emp-inc	0.558659	0.441341
Self-emp-not-inc	0.285714	0.714286
State-gov	0.268960	0.731040
Without-pay	0.000000	1.000000
All	0.248922	0.751078



Load the income data set and perform following operations


- 1.create a bar plot of Education VS salary status
- 2.create a cross table Education Vs salary status

```
#importing data
data=pd.read_csv("income.csv",na_values=[" ?"])
data2=data.dropna(axis=0)
sns.countplot(y="EdType",data=data2,hue="SalStat")
ed_tab=pd.crosstab(index=data2["EdType"],
                    columns=data2["SalStat"],
                    margins=True,normalize="index")
print(ed_tab)
```

The above table we can see that people who have done
Doctorate, Masters
, prof-schools are more likely to earn above
50000 USD per year when compared to others
Hence an influencing variable in avoiding the misuse of
subsidies



SalStat	greater than 50,000	less than or equal to 50,000
EdType		
10th	0.071951	0.928049
11th	0.056298	0.943702
12th	0.076923	0.923077
1st-4th	0.039735	0.960265
5th-6th	0.041667	0.958333
7th-8th	0.062837	0.937163
9th	0.054945	0.945055
Assoc-acdm	0.253968	0.746032
Assoc-voc	0.263198	0.736802
Bachelors	0.421491	0.578509
Doctorate	0.746667	0.253333
HS-grad	0.164329	0.835671
Masters	0.564229	0.435771
Preschool	0.000000	1.000000
Prof-school	0.749077	0.250923
Some-college	0.200060	0.799940
All	0.248922	0.751078



Load the income data set and perform following operations

- 1.create a bar plot of Occupation VS salary status
- 2.create a cross table occupation Vs salary status

```
#importing data
data=pd.read_csv("income.csv",na_values=[" ?"])
data2=data.dropna(axis=0)
sns.countplot(y="occupation",data=data2,hue="SalStat")
ocup_tab=pd.crosstab(index=data2["occupation"],
                      columns=data2["SalStat"],
                      margins=True,normalize="index")


print(ocup_tab)
```



SalStat	greater than 50,000	less than or equal to 50,000
occupation		
Adm-clerical	0.133835	0.866165
Armed-Forces	0.111111	0.888889
Craft-repair	0.225310	0.774690
Exec-managerial	0.485220	0.514780
Farming-fishing	0.116279	0.883721
Handlers-cleaners	0.061481	0.938519
Machine-op-inspct	0.124619	0.875381
Other-service	0.041096	0.958904
Priv-house-serv	0.006993	0.993007
Prof-specialty	0.448489	0.551511
Protective-serv	0.326087	0.673913
Sales	0.270647	0.729353
Tech-support	0.304825	0.695175
Transport-moving	0.202926	0.797074
All	0.248922	0.751078



Those who are making more than 50000 USD per year likely to work
as manager and professional,
hence important variable in avoiding misuse of subsidies

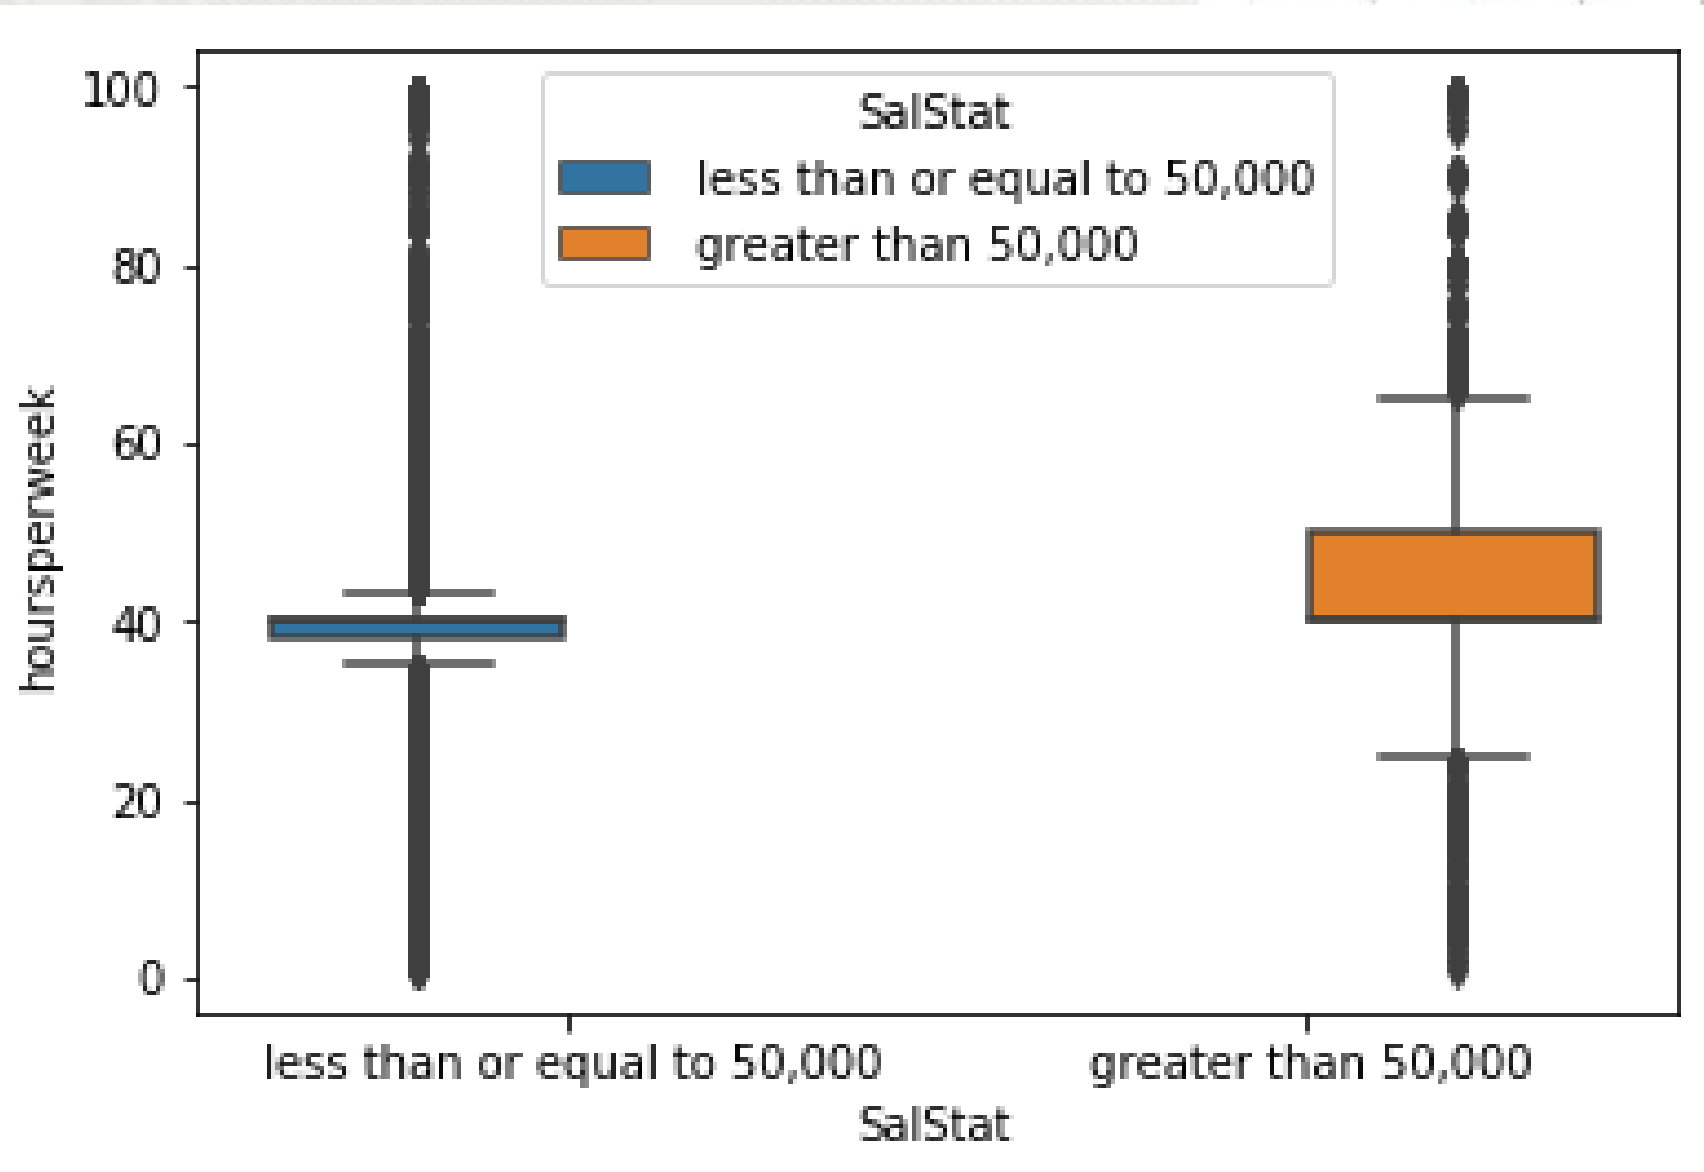


1.create a boxplot hourperweek Vs salary status

```
#importing data
data=pd.read_csv("income.csv",na_values=[" ?"])
data2=data.dropna(axis=0)
sns.boxplot(x=data2["SalStat"],
            y=data2["hoursperweek"],
            data=data2,hue="SalStat")
hrsweek_tab=pd.crosstab(index=data2["hoursperweek"],
                        columns=data2["SalStat"],
                        margins=True,normalize="index")
print(hrsweek_tab.head(10))
```



SalStat	greater than 50,000	less than or equal to 50,000
hoursperweek		
1	0.142857	0.857143
2	0.133333	0.866667
3	0.041667	0.958333
4	0.074074	0.925926
5	0.157895	0.842105
6	0.100000	0.900000
7	0.105263	0.894737
8	0.058824	0.941176
9	0.058824	0.941176
10	0.058559	0.941441





```
#Logistic Regression
"""
Data encoding
the salary status categories are encoded as 0,1
"""

#importing data
data=pd.read_csv("income.csv",na_values=[" ?"])
data2=data.dropna(axis=0)
#create a cross table on salstatus
print(data2["SalStat"].value_counts())
cat_salStat={" less than or equal to 50,000":0," greater than 50,000":1}
data2["SalStat"].replace(cat_salStat,inplace=True)
print(data2["SalStat"].head(6))
new_data=pd.get_dummies(data2,drop_first=True)
#storing the column names
column_list=list(new_data.columns)
print(column_list)
```




```
#seperating input features from the data
features=list(set(column_list)-set(["SalStat"]))
print(features)
#Storing output variable values in y
y=new_data["SalStat"].values
print(y)
#storing input feature values in x
x=new_data[features].values
print(x)
```



```
#splitting the data into train and test
train_x,test_x,train_y,test_y=train_test_split(x,y,test_size=0.3,
                                              random_state=0)

#creating a instance of logistict regression
logistic=LogisticRegression()
#fitting the values for x and y
logistic.fit(train_x,train_y)
#To display the fitting function attributes such as coef,intercept etc..
print(logistic.coef_)
print(logistic.intercept_)
```



```
#prediction from the test data
prediction=logistic.predict(test_x)
print(prediction)
#model evolution using classification metrics
#confusion metrics - To display correctly classified data
#and wrongly classified data
confus_matrix=confusion_matrix(test_y,prediction)
print(confus_matrix)
#Calculate the accuracy
accu_score=accuracy_score(test_y,prediction)
print(accu_score)
#To display missclassified values from the prediction
print("Missclassified")
print((test_y!=prediction).sum())
```



```
[0 0 0 ... 0 0 0]  
[[6338  485]  
 [ 941 1285]]  
0.8424135263565035  
Missclassified  
1426
```


Conclusion

You are aware of
Data Encoding
Project Life Cycle

We will proceed with
Case Study



**THANK
YOU**