



# SONET

## *Data Science* (Level-1)

### Machine Learning



# Data Preprocessing

# Steps for Machine Learning



Look at the big picture.

Get the data.

[Files/web/pandas](#)

Discover and visualize the data to gain insights.

Prepare the data for Machine Learning algorithms.

Select a model and train it.

Train the model

Fine-tune your model.

Present your solution.

Launch, monitor, and maintain your system



## Selection of Domain

Data Set: California census data

Metrics for each block group:

Population, median income, median housing price, ...

Problem:

Build a model to learn from this data and be able to predict the **median housing price** in any district, given median\_income

**As a Data Scientist , follow the steps of machine learning by starting with:**

Frame the Problem

Select the Performance Measure

Check the Assumptions

## Frame the Problem

Supervised/Unsupervised/reinforcement/classification/regressions/?

Batch Learning / Online Learning?

Supervised Learning – Labelled data is available

Regression – Asked to predict a value

Single regression problem – use median income to predict median\_house\_price

Batch learning is fine – small data, no rapid changes, no continues flow

## Select Performance Measure

There are various performance measures for each model.

We selected the regression model for our problem

Root Mean Square Error(RMSE)

$$\text{RMSE}(\mathbf{X}, h) = \sqrt{\frac{1}{m} \sum_{i=1}^m (h(\mathbf{x}^{(i)}) - y^{(i)})^2}$$

Mean Absolute Error (MAE)

$$\text{MAE}(\mathbf{X}, h) = \frac{1}{m} \sum_{i=1}^m |h(\mathbf{x}^{(i)}) - y^{(i)}|$$

m – Number of observations

h – Hypothesis function

$h(\mathbf{x}^{(i)})$  – Predicted Value for  $i^{\text{th}}$  instance

$y^{(i)}$  – Actual Value for  $i^{\text{th}}$  instance

RMSE is more sensitive to outliers than the MAE.



**Look at the big picture.**

## Check Assumptions

Make the assumptions that what they can expect in different way.

There is the chance of expectation of result as “cheap”, “reasonable”, “ Costly”

If so, there will be the wastage of working hours and efforts

Discuss with management and team and get clarity on this assumption

Lets continue with regression for now....

## Get Data

Downloading

Loading

Creating Training and test data

## Downloading

Data may be available in various forms/ sources

Web

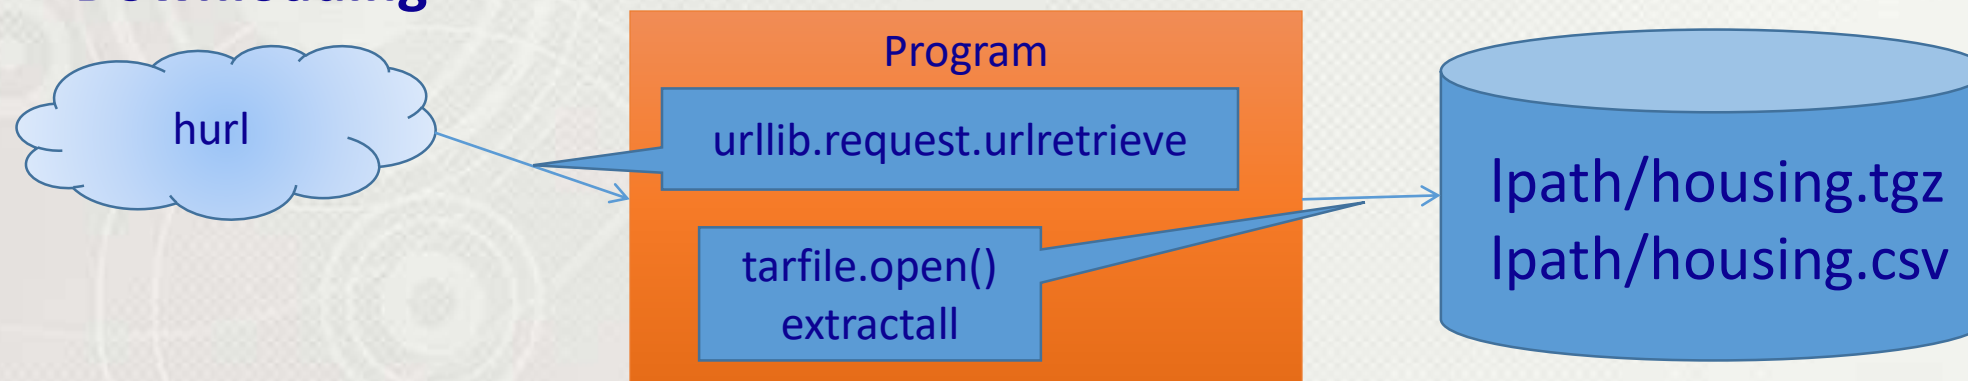
Files

Databases

Lets take CSV from given url.



### Downloading




```
import os
import tarfile
```

```
import urllib
```

```
hurl="https://raw.githubusercontent.com/ageron/handson-ml/master/datasets/housing/housing.tgz"
lpath="datasets/housing/"
urllib.request.urlretrieve(hurl,lpath+"housing.tgz")
htgz=tarfile.open(lpath+"housing.tgz")
htgz.extractall(path=lpath)
htgz.close()
```

Downloads the file from hurl into local path

## Loading Data into data frame (pandas)



```
import pandas as pd
d=pd.read_csv(lpath+"housing.csv")
d.head()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0	NEAR BAY
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0	NEAR BAY
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0	NEAR BAY

## Loading Data into data frame (pandas)

Let's see the various statistics of loaded dataframe

```
In [16]: d.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
longitude                20640 non-null float64
latitude                 20640 non-null float64
housing_median_age       20640 non-null float64
total_rooms               20640 non-null float64
total_bedrooms           20433 non-null float64
population                20640 non-null float64
households                20640 non-null float64
median_income             20640 non-null float64
median_house_value        20640 non-null float64
ocean_proximity           20640 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
In [19]: d.total_rooms.count()
```

```
Out[19]: 20640
```

```
In [17]: d['ocean_proximity'].value_counts()
```

```
Out[17]: <1H OCEAN      9136
INLAND           6551
NEAR OCEAN       2658
NEAR BAY         2290
ISLAND              5
Name: ocean_proximity, dtype: int64
```



## Loading Data into data frame (pandas)

Let's see the various statistics of loaded dataframe

```
In [22]: d.describe()
```

Out[22]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870671	206855.816909
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	382.329753	1.899822	115395.615874
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.499900	14999.000000
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.563400	119600.000000
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.534800	179700.000000
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.743250	264725.000000
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000100	500001.000000

```
In [24]: d.describe().population
```

Out[24]:

```
count    20640.000000
mean      1425.476744
std       1132.462122
min         3.000000
25%        787.000000
50%       1166.000000
75%       1725.000000
max       35682.000000
Name: population, dtype: float64
```

## Sample datasets in sklearn

<code>load_boston([return_X_y])</code>	Load and return the boston house-prices dataset (regression).
<code>load_iris([return_X_y])</code>	Load and return the iris dataset (classification).
<code>load_diabetes([return_X_y])</code>	Load and return the diabetes dataset (regression).
<code>load_digits([n_class, return_X_y])</code>	Load and return the digits dataset (classification).
<code>load_linnerud([return_X_y])</code>	Load and return the linnerud dataset (multivariate regression).
<code>load_wine([return_X_y])</code>	Load and return the wine dataset (classification).
<code>load_breast_cancer([return_X_y])</code>	Load and return the breast cancer wisconsin dataset (classification).

```
In [5]: import matplotlib.pyplot as plt
import numpy as np
from sklearn import datasets, linear_model
from sklearn.metrics import mean_squared_error, r2_score


iris=datasets.load_iris()
print(iris.DESCR)

type(iris.data)
```

```
Iris Plants Database
=====
```

matplotlib is a library to plot graphs and visualize data

```
In [36]: import matplotlib.pyplot as plt
d['median_house_value'].hist(bins=50,figsize=(10,5))
plt.show()
```





**Covered Initial two steps of Data Science Project**

**Look at Big Picture**

**Getting Data**

**Next:**

**Pre-Processing**

**THANK  
YOU**