

In [2]:

```
import pandas as pd
data = pd.read_csv("C:\\Users\\kmit\\Desktop\\housing.csv",",")
print(data.head(5))
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY

In [3]:

```
print("total samples....\n",data.size)
print("Null Values....\n", data.isnull().sum())
#print(data.isnull().count())
```

```
total samples....
206400
Null Values....
longitude          0
latitude           0
housing_median_age 0
total_rooms        0
total_bedrooms     207
population         0
households         0
median_income      0
median_house_value 0
ocean_proximity    0
dtype: int64
```

In [4]:

```
d1 = data.dropna(subset=['total_bedrooms'])
print(data.shape,d1.shape)
```

```
(20640, 10) (20433, 10)
```

In [5]:

d1.cov()

Out[5]:

	longitude	latitude	housing_median_age	total_rooms	total_b
<b>longitude</b>	4.014324	-3.957670	-2.758919	1.991284e+02	5.876
<b>latitude</b>	-3.957670	4.563981	0.320091	-1.711788e+02	-6.029
<b>housing_median_age</b>	-2.758919	0.320091	158.553558	-9.923225e+03	-1.700
<b>total_rooms</b>	199.128445	-171.178818	-9923.224538	4.775403e+06	8.567
<b>total_bedrooms</b>	58.768508	-60.299623	-1700.312817	8.567306e+05	1.775
<b>population</b>	227.660858	-263.874646	-4220.630517	2.122942e+06	4.191
<b>households</b>	43.286878	-58.619704	-1457.475788	7.677502e+05	1.578
<b>median_income</b>	-0.059174	-0.323087	-2.828672	8.213000e+02	-6.180
<b>median_house_value</b>	-10499.897668	-35669.333210	154703.602850	3.362452e+07	2.416

In [6]:

d1.corr()

Out[6]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
<b>longitude</b>	1.000000	-0.924616	-0.109357	0.045480	0.069608	0.100270	0.056513	-0.015550	-0.045398
<b>latitude</b>	-0.924616	1.000000	0.011899	-0.036667	-0.066983	-0.108997	-0.071774	-0.079626	-0.144638
<b>housing_median_age</b>	-0.109357	0.011899	1.000000	-0.360628	-0.320451	-0.295787	-0.302768	-0.118278	0.106432
<b>total_rooms</b>	0.045480	-0.036667	-0.360628	1.000000	0.930380	0.857281	0.918992	0.197882	0.133294
<b>total_bedrooms</b>	0.069608	-0.066983	-0.320451	0.930380	1.000000	0.877747	0.979728	-0.007723	0.049686
<b>population</b>	0.100270	-0.108997	-0.295787	0.857281	0.877747	1.000000	0.999999	0.197882	0.133294
<b>households</b>	0.056513	-0.071774	-0.302768	0.918992	0.979728	0.999999	1.000000	0.197882	0.133294
<b>median_income</b>	-0.015550	-0.079626	-0.118278	0.197882	-0.007723	0.197882	0.197882	1.000000	0.133294
<b>median_house_value</b>	-0.045398	-0.144638	0.106432	0.133294	0.049686	0.133294	0.133294	0.133294	1.000000

In [7]:

```
d1.corr()['median_house_value']
```

Out[7]:

```
longitude          -0.045398
latitude           -0.144638
housing_median_age  0.106432
total_rooms         0.133294
total_bedrooms      0.049686
population          -0.025300
households          0.064894
median_income       0.688355
median_house_value  1.000000
Name: median_house_value, dtype: float64
```

In [8]:

```
d1.corr()['median_house_value'].sort_values()[::-1]
```

Out[8]:

```
median_house_value  1.000000
median_income       0.688355
total_rooms         0.133294
housing_median_age  0.106432
households          0.064894
total_bedrooms      0.049686
population          -0.025300
longitude           -0.045398
latitude            -0.144638
Name: median_house_value, dtype: float64
```

In [9]:


```
corr_cols=d1.corr()['median_house_value'].sort_values()[::-1]
#corr_cols[1:4]
corr_cols.index
```

Out[9]:

```
Index(['median_house_value', 'median_income', 'total_rooms',
      'housing_median_age', 'households', 'total_bedrooms', 'population',
      'longitude', 'latitude'],
      dtype='object')
```


In [10]:

```
import matplotlib.pyplot as plt
plt.scatter(d1.median_income,d1.median_house_value)
plt.show()
```



In [11]:

```
plt.subplot(121)
plt.boxplot(d1.median_income)
plt.subplot(122)
plt.boxplot(d1.median_house_value)
plt.show()
```



In [12]:


```
d2=d1[d1.median_income<d1.median_income.quantile(0.8)]
d2.shape
```

Out[12]:

(16346, 10)


In [13]:

```
plt.subplot(121)
plt.boxplot(d2.median_income)
plt.subplot(122)
plt.boxplot(d2.median_house_value)
plt.show()
```




In [14]:

```
import matplotlib.pyplot as plt
plt.scatter(d2.median_income, d2.median_house_value)
plt.show()
```



In [15]:

```
import matplotlib.pyplot as plt
plt.scatter(d2.median_income,d2.median_house_value/10000)
plt.show()
```



In [32]:

```

d3=d2.drop(corr_cols.index[4:],axis=1)
d3=d3.drop(['ocean_proximity'],axis=1)
d3=d3.drop(['median_income'],axis=1)
print(d3.shape)
d3
d3.median_house_value = d3.median_house_value/100000
d3.total_rooms = d3.total_rooms/1000
d3.housing_median_age = d3.housing_median_age/10
print(d3)

```

(16346, 3)

	housing_median_age	total_rooms	median_house_value
4	5.2	1.627	3.422
5	5.2	0.919	2.697
6	5.2	2.535	2.992
7	5.2	3.104	2.414
8	4.2	2.555	2.267
9	5.2	3.549	2.611
10	5.2	2.202	2.815
11	5.2	3.503	2.418
12	5.2	2.491	2.135
13	5.2	0.696	1.913
14	5.2	2.643	1.592
15	5.0	1.120	1.400
16	5.2	1.966	1.525
17	5.2	1.228	1.555
18	5.0	2.239	1.587
19	5.2	1.503	1.629
20	4.0	0.751	1.475
21	4.2	1.639	1.598
22	5.2	2.436	1.139
23	5.2	1.688	0.997
24	5.2	2.224	1.326
25	4.1	0.535	1.075
26	4.9	1.130	0.938
27	5.2	1.898	1.055
28	5.0	2.082	1.089
29	5.2	0.729	1.320
30	4.9	1.916	1.223
31	5.2	2.153	1.152
32	4.8	1.922	1.104
33	4.9	1.655	1.049
...	...	...	...
20610	2.8	2.130	0.455
20611	2.7	1.783	0.470
20612	2.6	1.377	0.483
20613	3.1	1.728	0.534
20614	2.6	2.276	0.580
20615	2.3	1.076	0.575
20616	1.5	1.810	0.551
20617	2.0	0.561	0.708
20618	2.5	1.332	0.634
20619	2.2	1.891	0.991
20620	4.0	0.198	1.000
20621	3.7	1.244	0.775
20622	2.0	0.755	0.670
20623	3.2	1.158	0.655

9/7/2018

Multiple Linear Regression(Housing)

20625	3.7	0.102	0.720
20626	3.6	1.124	0.938
20627	0.5	0.358	1.625
20628	1.9	2.043	0.924
20629	2.8	10.035	1.083
20630	1.1	2.640	1.120
20631	1.5	2.655	1.072
20632	1.5	2.319	1.156
20633	2.7	2.080	0.983
20634	2.8	2.332	1.168
20635	2.5	1.665	0.781
20636	1.8	0.697	0.771
20637	1.7	2.254	0.923
20638	1.8	1.860	0.847
20639	1.6	2.785	0.894

[16346 rows x 3 columns]



In [39]:

```

import numpy as np
from sklearn import linear_model
from sklearn.metrics import mean_squared_error

#plt.scatter(d3["median_income"],d3["median_house_value"],color='b')

#print(sample1.head(1))
testsize=(int)(d3.shape[0]*0.30)

train=d3[:-testsize]
    #print(train.shape)
test=d3[-testsize:]
    #print(test.shape)

train_x=train[['housing_median_age','total_rooms']]
#train_x=train_x[:,np.newaxis]
train_y=train['median_house_value']
train_y=train_y[:,np.newaxis]
test_x = test[['housing_median_age','total_rooms']]
#test_x = test_x[:,np.newaxis]
test_y=test['median_house_value']
test_y=test_y[:,np.newaxis]
train_x.shape

print(train_x)

#print(".....",X_train_poly.shape,train_y.shape,X_test_poly.shape,test_y.shape)
lm1=linear_model.LinearRegression()
lm1.fit(train_x,train_y)
train_pred=lm1.predict(train_x)
#train_y-train_pred

#plt.scatter(train.median_income,train.median_house_value)
test_pred=lm1.predict(test_x)
    #plt.scatter(test_x,test_y)
#ax=plt.axis(projection='3d')

#ax.plot3D(test_x.total_rooms,test_x.housing_median_age, test.y)
plt.plot(test_x,test_y, 'g^')

plt.plot(test_x,test_pred,'r+')

print("MSE : ",mean_squared_error(test_y,test_pred), "\tRMSE:",np.sqrt(mean_squared_error(t
    #print(sample1.median_income.size,train.median_income.size)
plt.show()

```

	housing_median_age	total_rooms
4	5.2	1.627
5	5.2	0.919
6	5.2	2.535
7	5.2	3.104
8	4.2	2.555
9	5.2	3.549
10	5.2	2.202
11	5.2	3.503

13	5.2	0.696
14	5.2	2.643
15	5.0	1.120
16	5.2	1.966
17	5.2	1.228
18	5.0	2.239
19	5.2	1.503
20	4.0	0.751
21	4.2	1.639

In [26]:

```
import matplotlib.pyplot as plt
plt.boxplot(d3.total_bedrooms)
plt.show()
```

-----  
**AttributeError** Traceback (most recent call last)

<ipython-input-26-e1b635c37c1d> in <module>()

1 import matplotlib.pyplot as plt

----> 2 plt.boxplot(d3.total\_bedrooms)

3 plt.show()

~\Anaconda3\lib\site-packages\pandas\core\generic.py in \_\_getattr\_\_(self, name)

3079 if name in self.\_info\_axis:

3080 return self[name]

-> 3081 return object.\_\_getattribute\_\_(self, name)

3082

3083 def \_\_setattr\_\_(self, name, value):

**AttributeError:** 'DataFrame' object has no attribute 'total\_bedrooms'

In [ ]: